

EDA

The datasheet for 2009-10 and sheet containing data for 2010-11 are both same. So we will do EDA on one of the datasheet.

```
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from datetime import datetime
```

```
sns.set_style("whitegrid")
```

```
df1 = pd.read_csv("RetailPulseAssignmentData(2009-10).csv")
df2 = pd.read_csv("RetailPulseAssignmentData(2010-11).csv")
```

```
df1.head(4)
```

	Invoice	StockCode	Description	Quantity	\
0	489434	85048	15CM CHRISTMAS GLASS BALL 20 LIGHTS	12	
1	489434	79323P	PINK CHERRY LIGHTS	12	
2	489434	79323W	WHITE CHERRY LIGHTS	12	
3	489434	22041	RECORD FRAME 7" SINGLE SIZE	48	

	InvoiceDate	Price	Customer ID	Country
0	01-12-2009 07:45	6.95	13085.0	United Kingdom
1	01-12-2009 07:45	6.75	13085.0	United Kingdom
2	01-12-2009 07:45	6.75	13085.0	United Kingdom
3	01-12-2009 07:45	2.10	13085.0	United Kingdom

```
df1.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 525461 entries, 0 to 525460
Data columns (total 8 columns):
#   Column          Non-Null Count  Dtype
---  -
0   Invoice          525461 non-null object
1   StockCode       525461 non-null object
2   Description     522533 non-null object
3   Quantity        525461 non-null int64
4   InvoiceDate     525461 non-null object
5   Price           525461 non-null float64
6   Customer ID    417534 non-null float64
7   Country         525461 non-null object
dtypes: float64(2), int64(1), object(5)
memory usage: 32.1+ MB
```

```
df1.columns
```

```
Index(['Invoice', 'StockCode', 'Description', 'Quantity',
       'InvoiceDate',
```

```
    'Price', 'Customer ID', 'Country'],
    dtype='object')
```

```
df1[['Price', 'Customer ID']].apply(pd.to_numeric)
df1['Quantity'].astype('int64')
df1[['Description', 'Country']].astype(str)
print("The datatypes of columns are changed.")
```

The datatypes of columns are changed.

```
df1.describe()
```

	Quantity	Price	Customer ID
count	525461.000000	525461.000000	417534.000000
mean	10.337667	4.688834	15360.645478
std	107.424110	146.126914	1680.811316
min	-9600.000000	-53594.360000	12346.000000
25%	1.000000	1.250000	13983.000000
50%	3.000000	2.100000	15311.000000
75%	10.000000	4.210000	16799.000000
max	19152.000000	25111.090000	18287.000000

```
print("Does quantity have negative values:", (df1['Quantity'].values <
0).any())
```

Does quantity have negative values: True

```
neg_index = df1[df1['Quantity'].values < 0].index
df1.drop(neg_index,inplace = True)
```

The datasheet has **525461 transaction enteries** with records of

- (Invoice, StockCode, Description, Quality, InvoiceData, Price, CustomerID, Country)

```
print("Is the datasheet with 2009-10 and 2010-11 are exactly same:
",df1.equals(df2))
```

Is the datasheet with 2009-10 and 2010-11 are exactly same: False

```
print("Null values/missing values in the datasheet:-")
df1.isnull().sum()
```

Null values/missing values in the datasheet:-

Invoice	0
StockCode	0
Description	1101
Quantity	0
InvoiceDate	0
Price	0
Customer ID	105440
Country	0
dtype:	int64

```
df1.dropna(inplace=True) #Description and CustomerID cannot be
replaced, thus those enteries are dropped
df1.isnull().sum()
```

```
Invoice      0
StockCode    0
Description  0
Quantity     0
InvoiceDate  0
Price        0
Customer ID  0
Country      0
dtype: int64
```

```
df1 = df1.drop_duplicates()
```

```
df1['InvoiceDate'] = pd.to_datetime(df1['InvoiceDate'])
df1['Date'] = df1['InvoiceDate'].dt.strftime('%d-%m-%Y')
df1['Date'] = pd.to_datetime(df1['Date'])
df1['Time'] = df1['InvoiceDate'].dt.strftime('%H:%M:%S')
```

```
df1['TotalPrice'] = df1['Quantity'] * df1['Price']
```

```
df1.head(4)
```

	Invoice	StockCode	Description	Quantity	\
0	489434	85048	15CM CHRISTMAS GLASS BALL 20 LIGHTS	12	
1	489434	79323P	PINK CHERRY LIGHTS	12	
2	489434	79323W	WHITE CHERRY LIGHTS	12	
3	489434	22041	RECORD FRAME 7" SINGLE SIZE	48	

	InvoiceDate	Price	Customer ID	Country	
0	2009-01-12 07:45:00	6.95	13085.0	United Kingdom	2009-12-01
1	2009-01-12 07:45:00	6.75	13085.0	United Kingdom	2009-12-01
2	2009-01-12 07:45:00	6.75	13085.0	United Kingdom	2009-12-01
3	2009-01-12 07:45:00	2.10	13085.0	United Kingdom	2009-12-01

	Time	TotalPrice
0	07:45:00	83.4
1	07:45:00	81.0
2	07:45:00	81.0
3	07:45:00	100.8

Analysing data

```
print(f"The range of dates of transactions: {df1['Date'].min()} to
{df1['Date'].max()}")
```

The range of dates of transactions: 2009-12-01 00:00:00 to 2010-12-09 00:00:00

Above shows the range of dates of which the data is available.

Now, These are the **frequency of transactions** done by the same Customer and what amount sums their purchases. The top 5 Customers who have been the most loyal Customers.

```
Freq_df = df1.groupby('Customer ID').agg({'Customer ID': lambda x:
len(x),
                                         'TotalPrice': lambda x:
x.sum(),
                                         'Date': lambda x: (x.max() -
x.min()).days})
```

```
Freq_df.rename(columns = {'Customer ID': 'Frequency of transactions'},
inplace= True)
```

```
Freq_df.sort_values(by = ['Frequency of
transactions'], ascending=False, inplace = True)
Freq_df['Customer ID'] = Freq_df.index
Freq_df['Customer ID'] = Freq_df['Customer ID'].astype(str)
Freq_df.rename(columns = {'Date': 'Period of purchases'}, inplace =
True)
```

```
print("The top five customer with most transaction: -")
top_five_customer_data = Freq_df.head(5)
top_five_customer_data
```

The top five customer with most transaction: -

purchases \ Customer ID	Frequency of transactions	TotalPrice	Period of
----------------------------	---------------------------	------------	-----------

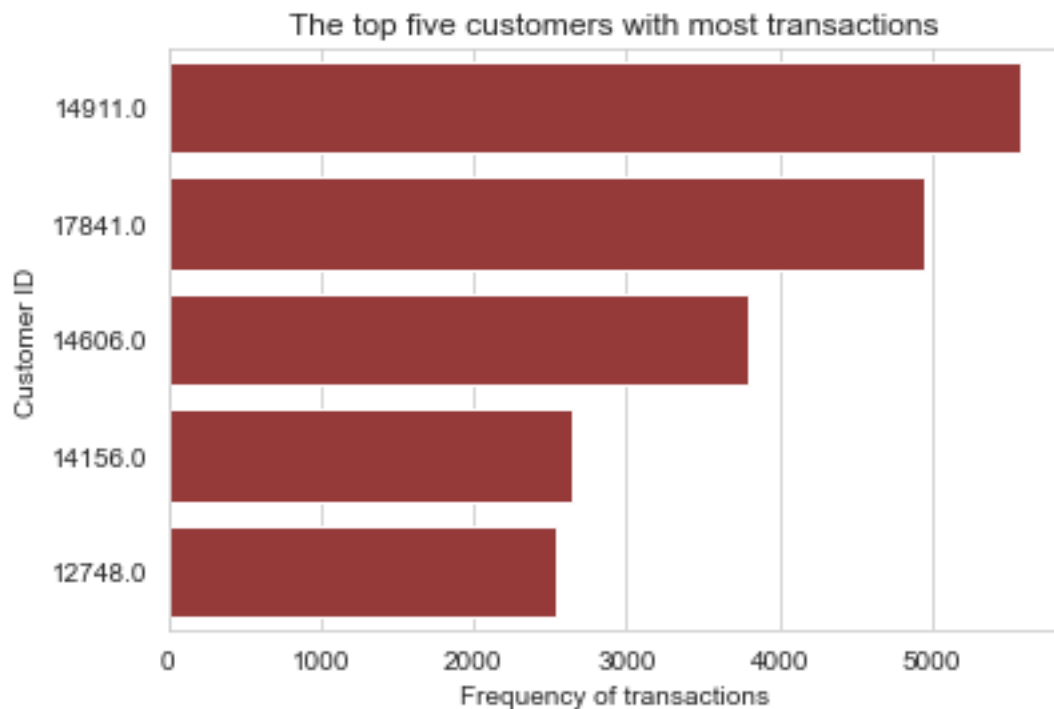
14911.0 373	5568.0	152121.22	
17841.0 372	4944.0	29562.02	
14606.0 370	3802.0	18482.10	
14156.0 367	2646.0	196549.74	
12748.0 370	2532.0	22457.90	

Customer ID	Customer ID
14911.0	14911.0
17841.0	17841.0

```
14606.0      14606.0
14156.0      14156.0
12748.0      12748.0
```

```
plt.title("The top five customers with most transactions")
sns.barplot(data = top_five_customer_data, x= 'Frequency of
transactions', y = 'Customer ID', color='brown')
```

```
<AxesSubplot:title={'center':'The top five customers with most
transactions'}, xlabel='Frequency of transactions', ylabel='Customer
ID'>
```



The loyal customers or it's value might be evaluated by repeated transactions of valuable products.

We can see the **Customer ID : 14911.0** had most frequent transactions resulting of **5568** purchases. followed by these four customers :-

1. ID :17841.0
2. ID :14606.0
3. ID :14156.0
4. ID :12748.0

*# Customer Value can be evaluated as: Average Order Value * Purchase Frequency*

```
Freq_df['avg_order_value'] = Freq_df['TotalPrice'] / Freq_df['Period
of purchases']
```

```

Freq_df['profit_margin'] = Freq_df['TotalPrice'] * 0.05
Freq_df['Customer_value'] = Freq_df['avg_order_value'] *
Freq_df['Frequency of transactions']
Freq_df['Customer_lifetime_value'] = Freq_df['Customer_value'] *
Freq_df['profit_margin']

```

```

Freq_df.sort_values(by = 'Customer_value',ascending = False)
Freq_df.head(5)

```

purchases \ Customer ID	Frequency of transactions	TotalPrice	Period of
14911.0 373	5568.0	152121.22	
17841.0 372	4944.0	29562.02	
14606.0 370	3802.0	18482.10	
14156.0 367	2646.0	196549.74	
12748.0 370	2532.0	22457.90	

Customer ID	avg_order_value	profit_margin
14911.0 2.270807e+06	407.831689	7606.061
17841.0 3.928888e+05	79.467796	1478.101
14606.0 1.899161e+05	49.951622	924.105
14156.0 1.417086e+06	535.557875	9827.487
12748.0 1.536849e+05	60.697027	1122.895

Customer ID	Customer_lifetime_value
14911.0	1.727190e+10
17841.0	5.807293e+08
14606.0	1.755024e+08
14156.0	1.392640e+10
12748.0	1.725720e+08

We can see the similar results where **Customer ID : 14911.0** had still **most valued** along with other 4 people who are in the top five customers.

```
df1['month'] = pd.DatetimeIndex(df1['Date']).month

date_df = df1.groupby('Date').agg({'TotalPrice': lambda x: x.sum(),
                                   'Quantity': lambda x: x.sum()})
```

```
date_df['date'] = date_df.index
date_df['date'] = pd.to_datetime(date_df['date'])
date_df['month'] = date_df['date'].dt.month
# The months are added as column
```

```
date_price_df = date_df.sort_values(by = ['TotalPrice'], ascending =
False)
print("The highest transaction amount recorded on a day : ",
date_price_df['TotalPrice'][0])
date_price_df.head(4)
```

The highest transaction amount recorded on a day : 115172.23999999999

Date	TotalPrice	Quantity	date	month
2010-09-27	115172.24	125369	2010-09-27	9
2010-10-14	86945.13	37130	2010-10-14	10
2010-11-04	83876.12	58902	2010-11-04	11
2010-11-10	73575.93	33426	2010-11-10	11

```
date_quantity_df = date_df.sort_values(by = ['Quantity'], ascending =
False)
print("The highest units purchased on a day :
",date_price_df['Quantity'][0])
date_quantity_df.head(4)
```

The highest units purchased on a day : 125369


Date	TotalPrice	Quantity	date	month
2010-09-27	115172.240	125369	2010-09-27	9
2010-08-09	31336.590	100084	2010-08-09	8
2010-02-15	42318.211	96358	2010-02-15	2
2010-03-17	29543.780	78565	2010-03-17	3

We can see that on **27-09-2010** the highest transaction both in **quantity and amount / price** was recorded globally. The quantity sold was **125369 units**, all of which summed to an inflow of **115172.240 USD** (currency it was recorded).

```
plt.figure(figsize = (15,5))
```

```
date_df.sort_values(by = 'month', ascending = True, inplace = True)
sns.barplot(data = date_df, x = 'month', y = 'TotalPrice')
```


```
<AxesSubplot:xlabel='month', ylabel='TotalPrice'>
```



Most revenue generated with transactions were made in month of **November** followed by Octobe and December. Thus this likely indicates the most inflow occurs in winter season which is the end of the year.

```
plt.figure(figsize = (15,5))
Customer_Country = df1.groupby('Country')['Customer
ID'].nunique().sort_values(ascending=False).reset_index().head(11)
sns.barplot(data=Customer_Country,x='Country',y='Customer
ID',palette='Blues')
```

<AxesSubplot:xlabel='Country', ylabel='Customer ID'>



Almost all the Customers are residing in *United Kindom*.

```
df1['Description'].unique()
array(['15CM CHRISTMAS GLASS BALL 20 LIGHTS', 'PINK CHERRY LIGHTS',
      ' WHITE CHERRY LIGHTS', ..., 'BAKING MOULD ROSE MILK
CHOCOLATE',
      'BAKING MOULD CHOCOLATE CUP CAKES',
      'BAKING MOULD EASTER EGG MILK CHOC'], dtype=object)
len(df1['Description'].unique())
4444
```



```
Desc_df = df1.groupby('Description').agg({'Customer ID': lambda x:
len(x),
'TotalPrice': lambda x:
x.sum(),
'Quantity': lambda x:
x.sum()})
```

```
Desc_df_ID = Desc_df.sort_values(by = 'Customer ID', ascending =
False)
Desc_df_price = Desc_df.sort_values(by = 'TotalPrice', ascending =
False)
Desc_df_Quantity = Desc_df.sort_values(by = 'Quantity',ascending =
False)
```

```
Desc_df_ID.head(5)
```

Description	Customer ID	TotalPrice	Quantity
WHITE HANGING HEART T-LIGHT HOLDER	3107.0	151339.16	56814
REGENCY CAKESTAND 3 TIER	1696.0	143727.60	12489
STRAWBERRY CERAMIC TRINKET BOX	1372.0	29694.86	25330
ASSORTED COLOUR BIRD ORNAMENT	1355.0	70291.03	44431
HOME BUILDING BLOCK WORD	1192.0	28028.30	4784


We can see that **"WHITE HANGING HEART T-LIGHT HOLDER"** is the best seller followed by other items in above dataframe.

There were **4444 different items/products** that were purchased. So people preferring to buy same things can be assumed as a category of people preferring that item. The most preferred product was WHITE HANGING HEART T-LIGHT HOLDER with over 56814 units sold contributing in **56814 USD** of inflow.

```
plt.figure(figsize = (30,10))
plt.title("The units sold with the overall amount / Price")
```

```
sns.lineplot(x = date_df['date'], y = date_df['TotalPrice'])
sns.lineplot(x = date_df['date'], y = date_df['Quantity'], color =
'black')
plt.legend(['TotalPrice', 'Quantity'])
```

```
<matplotlib.legend.Legend at 0x22f4475ceb0>
```



Here we can see the most of **units were sold around october and november** and most inflow was around the same time. With more data seasonality could have been detected confirming if there is some periodic pattern of purchase but with this data we can't confirm the same.